

INNOVATION

DATABASE ADMINISTRATION

TI. Database Server Architecture

C3: Oracle Storage Structures

Plan

- 2. Operating Oracle Database Server Internals
 - 2.1 Structures of Oracle Instance
 - Server Processes, Background Processes, Internal Memory
 - 2.2 Oracle Instance life cycle
 - Starting and stopping stages, Oracle Instance Data Query Processing
 - 2.3 Oracle Database Storage Structures
 - Data Files, Tablespaces, Blocks and extents, Data Segments
 - Database creation process
 - 2.4 Table storage
 - Table space allocation, Table segments, Cluster Segments, Partitions
 - 2.5 Index storage
 - Index organization, Index space allocation, IO-Tables, Index Partitions

2.4 Storing SQL Tables

Table storage structures

- Data Storage Types/Data Segment Types for SQL Data in tables and indexes:
 - Table Data Storage:
 - TABLE segments (conventional data storage);
 - CLUSTER segments;
 - TABLE PARTITION segments;
 - INDEX segments;
 - Index Data Storage:
 - INDEX segments;
 - INDEX PARTITION segments.

CLUSTER Segments to store SQL table rows

- Rows from a clustered table will be physically grouped by a cluster-key within table physical storage.
- Rows from different logical tables (usually related by a foreign key - parent key relationship) could be located the same table-block by sharing the same grouping key.
- The rows of clustered tables could be grouped:
 - by some indexing (and sorting) key resulting an *index cluster*;
 - by some hash-based key resulting a *hash cluster*: the same *hash function* will be used:
 - at **INSERT-time** to determine physical location (data block id) to store the new rows and_
 - at **SELECT-time** to access physical rows by filtering predicates based on the cluster key.

INDEX and HASH Cluster CREATE commands

```
CREATE CLUSTER cluster
( column datatype
  [ column datatype ] ...)
[PCTFREE integer]
[PCTUSED integer]
[INITRANS integer]
[ MAXTRANS integer]
[ SIZE integer [ K|M ]
  [ STORAGE storage-clause]
  [ LOGGING | NOLOGGING]
  [ NOSORT]
  [ TABLESPACE tablespace]
  [ INDEX];
CREATE INDEX ... ON CLUSTER ...;
CREATE TABLE ... CLUSTER ...(key);
```

```
CREATE CLUSTER cluster
( column datatype
  [ column datatype ] ...)
HASHKEYS integer
[SINGLE TABLE]
[ HASH IS expression]
[PCTFREE integer]
[PCTUSED integer]
[INITRANS integer]
[ MAXTRANS integer]
[ SIZE integer [ K|M ]
  [ STORAGE storage-clause]
  [ LOGGING | NOLOGGING]
  [ NOSORT]
  [ TABLESPACE tablespace];
CREATE TABLE ... CLUSTER ...(key);
```


Dictionary Views on Clusters

DBA_CLUSTERS

OWNER
CLUSTER_NAME
TABLESPACE_NAME
KEY_SIZE
CLUSTER_TYPE
FUNCTION
HASHKEYS

DBA_TAB_COLUMNS

OWNER
TABLE_NAME
COLUMN_NAME
DATA_TYPE
DATA_LENGTH
DATA_PRECISION
DATA_SCALE

DBA_CLU_COLUMNS

OWNER
CLUSTER_NAME
CLU_COLUMN_NAME
TABLE_NAME
TAB_COLUMN_NAME

DBA_CLUSTER_HASH_EXPRESSIONS

OWNER
CLUSTER_NAME
HASH_EXPRESSION

IOT: Index Organized Tables

- The ROWs within IOTs will be entirely *ordered* by using a sort-key that will finally define the B*Tree *index* segment structure.
- An IOT has the same root-to-branches-to-leafs structure as a regular index, the difference consists in storing all row values and not just the ROWID into the leaf entries.

```
CREATE TABLE table
(column datatype [,column datatype ] ...)
ORGANIZATION INDEX
[TABLESPACE tablespace]
[PCTFREE integer]
[INITRANS integer]
[      [MAXTRANS integer]
      [STORAGE storage-clause]
      [PCTTHRESHOLD integer]
[INCLUDING column]]
[OVERFLOW segment_attributes_clause ]
```

Dictionary Views on IOTs

DBA_TABLES

**OWNER
TABLE_NAME
TABLESPACE_NAME
IOT_TYPE
IOT_NAME**

DBA_INDEXES

**OWNER
TABLE_NAME
INDEX_NAME
INDEX_TYPE
PCT_THRESHOLD
INCLUDE_COLUMN**



TABLE PARTITIONS

- **Splitting** tables into **partitions**:
 - assumes a grouping strategy for tables rows based on:
 - **range partitioning or interval partitioning**;
 - **list partitioning**, where partition key values will explicitly be declared;
 - **hash partitioning**, using hash functions on partition-keys to determine partition location for table rows.
 - produces a **storage distribution** for a single table as a set of data segments (of PARTITION type); those segments could eventually be distributed into different cluster nodes (by using TABLESPACE clause and TABLESPACE flexible mechanism to manage storage allocation).

RANGE PARTITION TYPE

```
CREATE TABLE tbl_range_partitioned (  
    range_key_col INTEGER,  
    misc_data_col VARCHAR2(100)  
)  
PARTITION BY RANGE (range_key_col)  
(  
    PARTITION ptbl_1 VALUES LESS THAN (1000)  
        TABLESPACE TS1,  
    PARTITION ptbl_2 VALUES LESS THAN (2000)  
        TABLESPACE TS2,  
    PARTITION ptbl_3 VALUES LESS THAN (3000)  
        TABLESPACE TS3,  
    PARTITION ptbl_4 VALUES LESS THAN (maxvalue)  
        TABLESPACE TS4  
);
```

Practice C3_P1

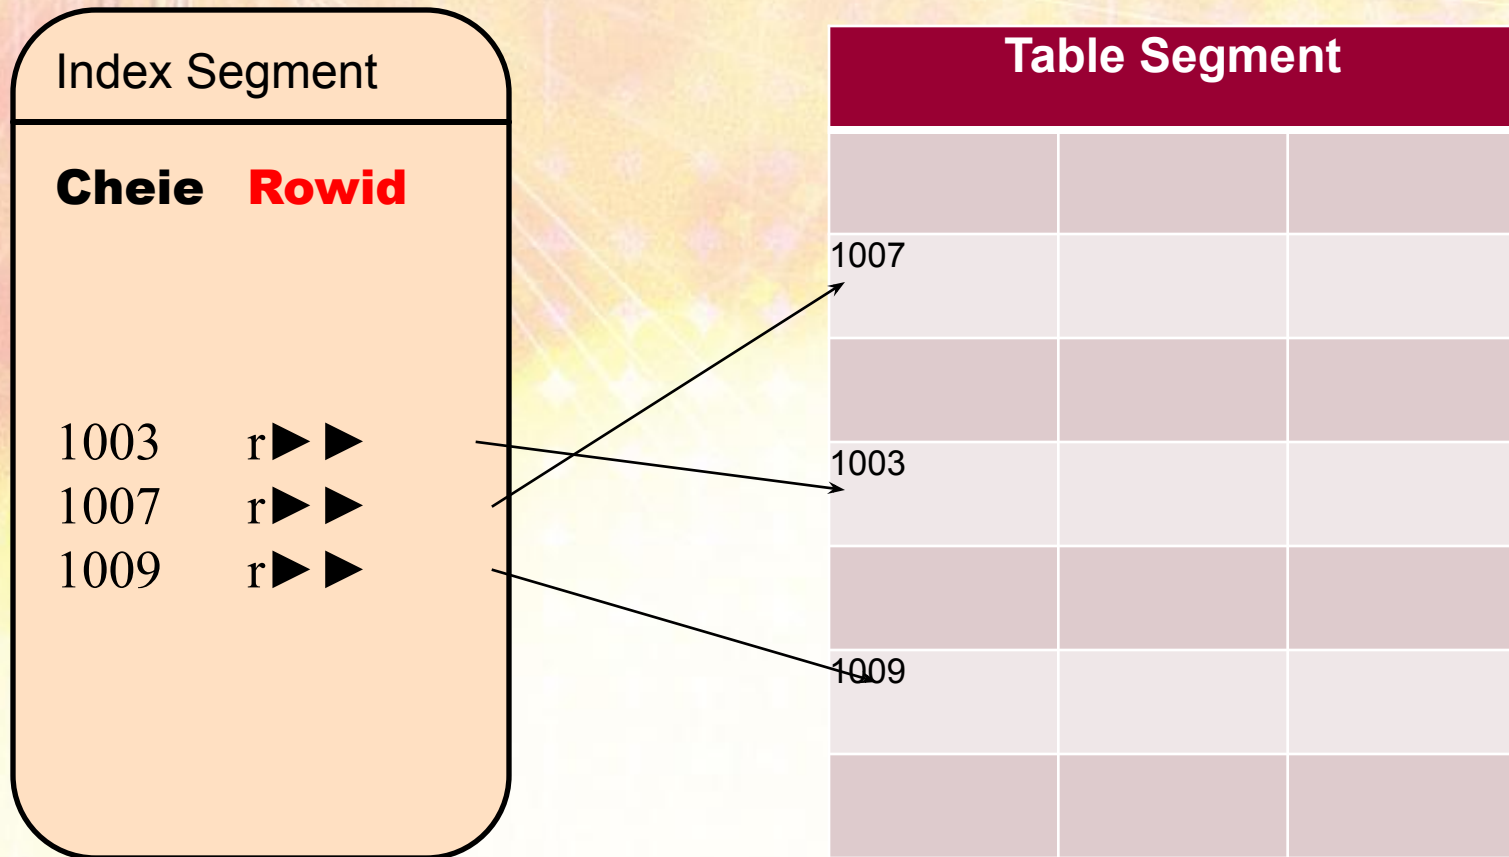
Script	Topics
C3_P1.1.Base_Schema_CREATE_TABLE_STORAGE_TYPES.sql	Recreate tables for OFA with different storage types: CREATE TABLE INDEX ORGANIZED CREATE CLUSTER CREATE PARTITION
C3_P1.2.TABLE_STORAGE_Stats.sql	Stats

2.5 Storing Indexes

INDEX Storage Structures

- **Indexes** represent the main **access paths** to table data:
 - index blocks are loaded before table blocks;
 - then index blocks are scanned by using SQL phrase **predicates** (WHERE clauses) to get the entries with the ROWIDs targeting the original table blocks.
- **Index internal organization:**
 - (unlike “heap” tables) each new entry will have a pre-computed location: index entries will be **physically ordered by the index key**;
 - index blocks are hierarchically organized in order to efficiently find/locate index entries containing the key-values coming matching the access predicates (from SELECT-WHERE clauses).

Index Role



ORACLE Index Classification

- 1. The logical perspective:
 - *number of columns criteria*:
 - single-column indexes;
 - composite indexes;
 - *uniqueness criteria*:
 - uniques indexes;
 - non-uniques indexes.
- 2. The internal implementation (physical) perspective:
 - B-Tree and bitmap index types;
 - Partitioned and non-partitioned indexes.

ORACLE B*TREE and Bitmap Indexes

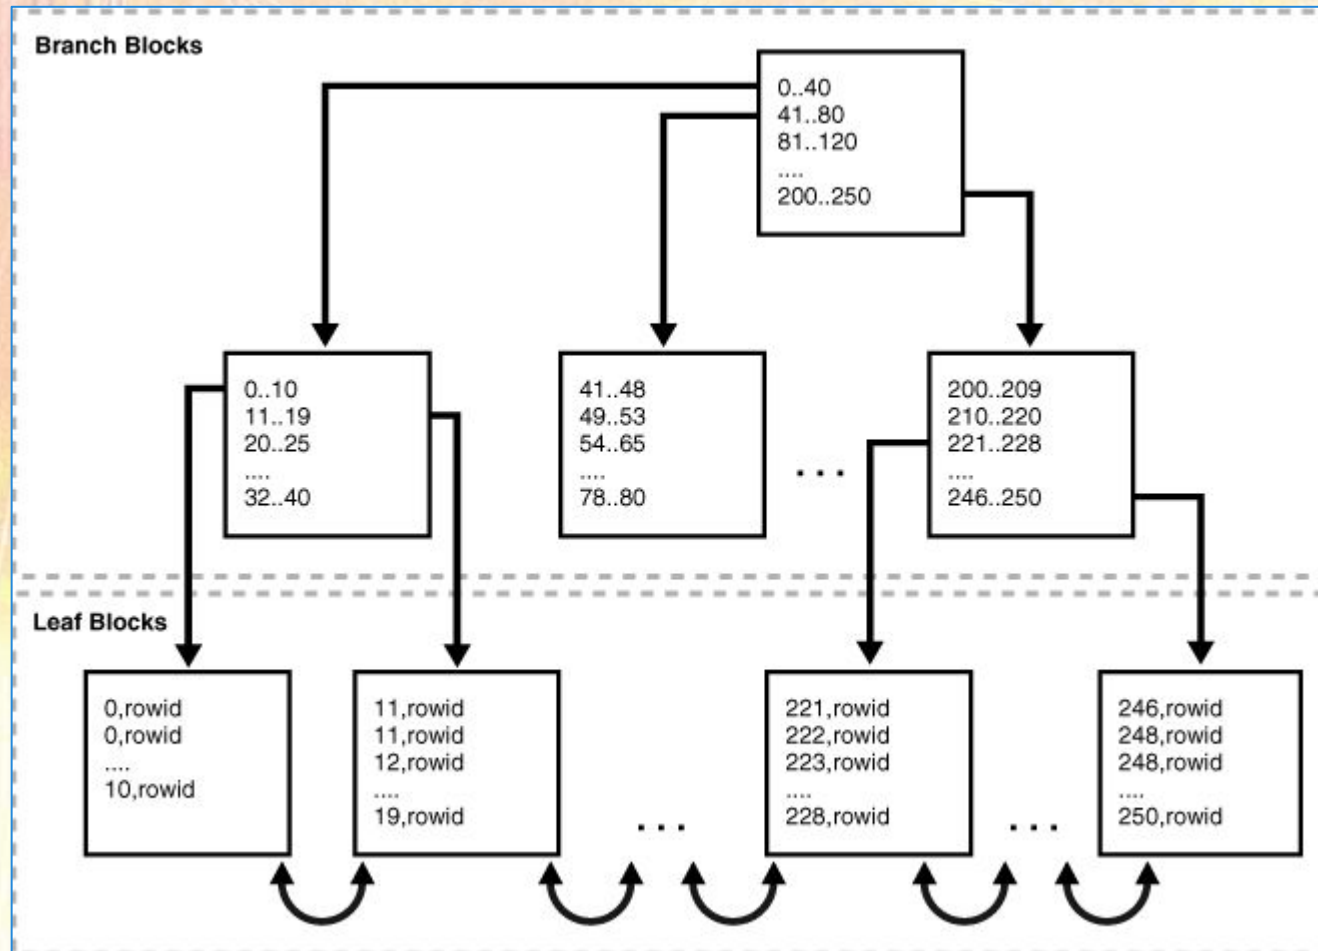
- The index Structure (B*tree or bitmap) relies on:
 - a specific inter-connected system of branch blocks:
 - starting from a unique root node containing a set of *pointers* to a **chained** set of branch nodes →
 - forming equally sized **search paths** with the same number of levels to scan; any key-value could start the search operation;
 - an **ordered set of leaf-nodes** containing associations of index-key-value and table ROWIDs.

ORACLE B*TREE and Bitmap Indexes


Leaf Nodes

- Leaf entries of **B*TREE indexes**:
 - associate each index-key-value to a single-ROWID (uniques indexes) or with a ROWID-set (non-unique indexes), each ROWID contains the location information (the file and table block) where the table row is stored.
- Leaf entries of **BITMAP indexes**:
 - associate each index-key-value to a BITMAP (set of bites) where every bit will correspond to a single table ROW:
 - if the bit is set to 1: the corresponding row contains the key-value;
 - if the bit is set to 0: the corresponding row do not contains the key-value.

B-Tree Index Structure



Oracle BITMAP Index



CREATE INDEX and ALTER INDEX Commands

```
CREATE [UNIQUE] INDEX
index ON table
( column [ column ] ...)
[TABLESPACE tablespace]
[PCTFREE integer]
[INITRANS integer]
[MAXTRANS integer]
[STORAGE storage-clause]
[LOGGING | NOLOGGING]
[NOSORT]
```

```
ALTER INDEX index REBUILD
TABLESPACE tablespacename
[PCTFREE integer]
[INITRANS integer]
[MAXTRANS integer]
[STORAGE storage-clause]
[LOGGING | NOLOGGING]
[NOSORT]
[ONLINE];
```

Dictionary Views on Indexes

DBA_INDEXES

OWNER
INDEX_NAME
INDEX_TYPE
TABLE_OWNER
TABLE_NAME
UNIQUENESS
TABLESPACE_NAME
LOGGING
STATUS



DBA_IND_COLUMNS

INDEX_OWNER
INDEX_NAME
TABLE_OWNER
TABLE_NAME
COLUMN_NAME
COLUMN_POSITION
COLUMN_LENGTH

INDEX PARTITIONS

- **Unpartitioned indexes:**
 - will be created by default for unpartitioned tables;
 - could be created for partitioned tables ignoring the actual table partitioning policy.
- **Partitioned Indexes** come from a partitioning policy:
 - specified at the table level and reflected at the index level in order to locate in the same storage space both: each index partition is associated with one table partition (**local indexes**);
 - specified at the index level: independent from table partitioning rule (**global indexes**).

GLOBAL INDEXES

- Index partitioning clause (PARTITION BY) not related with table partitioning clause:

```
CREATE TABLE tbl_part (t_id INTEGER, t_date DATE)
  PARTITION BY RANGE (t_date)
  ( PARTITION ptbl_1
    VALUES LESS THAN (TO_DATE('01/01/2019', 'DD/MM/YYYY')),
    PARTITION ptbl_2
    VALUES LESS THAN (TO_DATE('01/01/2020', 'DD/MM/YYYY'))
    PARTITION ptbl_3 VALUES LESS THAN (MAXVALUE));
```

```
CREATE INDEX tbl_idx_g ON tbl_part (t_date) GLOBAL
  PARTITION BY RANGE (t_date)
  ( PARTITION pidx_1
    VALUES LESS THAN (TO_DATE('15/10/2019', 'DD/MM/YYYY')),
    PARTITION pidx_2
    VALUES LESS THAN (TO_DATE('15/10/2020', 'DD/MM/YYYY'))
    PARTITION pidx_3 VALUES LESS THAN (MAXVALUE));
```


LOCAL Indexes

- Without using any specific index partitioning clause, the actual partitioning rule for the indexes of a partitioned table will depend on the PARTITION table policy:

```
CREATE TABLE tbl_part (t_id INTEGER, t_date DATE)
  PARTITION BY RANGE (t_date)
  ( PARTITION ptbl_1 VALUES LESS THAN TO_DATE('01/08/2019',
    'DD/MM/YYYY'),
    PARTITION ptbl_2 VALUES LESS THAN TO_DATE('01/01/2020',
    'DD/MM/YYYY'));
```

```
CREATE INDEX tbl_idx_loc ON tbl_part(t_date) LOCAL;
```


Practice C3_P2

Script	Topics
C3_P1.1.Base_Schema_CREATE_TABLE_STORAGE_TYPES.sql C1_P2.3.Base_Schema_INSERT.sql	Recreate indexes for OFA tables: CREATE INDEX CREATE INDEX LOCAL
C3_P2.1.INDEX_STORAGE_REBUILD.sql	ALTER INDEX REBUILD
C3_P2.2.INDEX_STORAGE_Stats.sql	Index stats

References

Titles

Craig S. Mullins, *Database Administration: the complete guide to practices and procedures*, Second Edition, Addison-Wesley, 2013

Thomas Kyte and Darl Kuhn, *Expert Oracle Database Architecture*, Third Edition, Apress, 2015

Lahdenmaki, Tapio, Leach, Michael, *Relational database index design and optimizers: DB2, Oracle, SQL server et al*, John Wiley & Sons, 2005

Bob Bryla, Kevin Loney *Oracle Database 11g DBA Handbook*, (Oracle Press), McGraw-Hill Osborne Media, 2008