

# DATABASE ADMINISTRATION

## T2: DATABASE OPTIMIZATION [C5]

Table Tuning and Index-based Optimization

# CHAPTER PLAN

- 1. Physical Design and Tools
  - 1.1 Physical Design - Optimization Objectives
  - 1.2 Physical Design Process
  - 1.3 Execution Plan Tool
  - 1.4 CBO Scopes
- **2. Table and Index Access Optimization**
  - **2.1 Buffering & Table (scan) Access Opt. Techniques**
  - **2.2 Index-based Optimization Techniques**
- 3. JOIN and SORT Operation Optimization
  - 3.1 JOIN Optimization Techniques
  - 3.2 SORT Optimization Techniques
- 4. More advanced techniques
  - Transformed Subqueries: subquery refactoring
  - Subquery optimization: temporary tables and materialized views
  - Access optimization with Parallel Query

## 2. TABLE and INDEX-based Optimization

- TABLE Storage and Access Tuning Techniques
- INDEX based Optimization:
  - to access tables;
  - to join tables;
  - to sort rows.

## 2.1 TABLE Optimization Techniques

- Table buffering
- Table STORAGE and table access paths OPERATIONS:
  - Full table scan;
  - Table access by index;
  - Table access by rowid;
  - Table access by hash;
  - Table access by cluster.

Storage Type	Access Path
HEAP	FTS, TABLE ACCESS BY INDEX ROWID
IOT	INDEX SCAN
CLUSTER HASH [SINGLE TABLE]	TABLE ACCESS BY HASH
CLUSTER INDEX [MULTIPLE TABLES]	TABLE ACCESS BY CLUSTER

- Table Tuning TECHNIQUES

# SQL Table Storage and Table Access Paths

- The physical design decision to choose the table storage type could be determined by some specific table access **strategies**:
  - *S1: Unordered table + primary key index + secondary indexes:*
    - Unordered tables could be stored as:
      - *Heap* storage structures;
      - *Hash* storage structures;
    - Indexing criteria could be determined by:
      - the initial primary and unique key constraints;
      - the columns coming from *equals* or *range* predicates within WHERE clauses of SELECT queries;
  - *S2: Ordered table + secondary indexes:*
    - Tables could be:
      - partially ordered: by using CLUSTER-like storage;
      - totally ordered: by using primary key index storage.

# SQL Table Storage and Table Access Paths

## When to use: heap, hash, cluster or index storage types

- When to use **HEAP** (unordered) storage:
  - massive data workloads occurs immediately after creation or at scheduled points in time;
  - the final table size is relatively small;
  - the access probability is the same for each and every table row;
  - most frequently access requests could be easily covered by additional indexes.
- When to use **HASH** (unordered) storage:
  - the access requests are very stable and are based on equals-like predicates;
  - hash-result values produced by hash-functions are highly controllable to accurately predict memory-page count number (the number of data blocks);
- When to use **CLUSTER-ordered** storage:
  - to pre-determine aggregation of row groups for aggregation queries;
  - to pre-determine JOIN-ing row groups (physical denormalization instead of logical denormalization) of related tables;
- When to use **INDEX-ordered** storage:
  - primary-key index is the main access criteria;
  - to pre-determine row ordering on final resultset.

# PL/SQL Hash Function Example

```
CREATE OR REPLACE
FUNCTION Get_Hash_Value(key_value VARCHAR) RETURN Number IS
  l_sum number := 0;
  l_tmp number := 0;
begin
  FOR i in 1..length(key_value) LOOP
    l_tmp := ascii(substr(key_value,i,1));
    l_sum := l_sum + l_tmp;
  END LOOP;
  return mod(l_sum,13);
END;
/
```

# Table Scan Operations



- Table access methods (OEP-operations):
  - table scan;
    - TABLE ACCESS FULL;
    - TABLE ACCESS SAMPLE;
  - access by rowid:
    - TABLE ACCESS BY USER ROWID;
    - TABLE ACCESS BY RANGE ROWID;
  - access by index:
    - TABLE ACCESS BY INDEX ROWID;
  - partitioning access for TABLE ACCESS FULL:
    - PARTITION RANGE SINGLE
    - PARTITION RANGE ITERATOR
  - full tables scan with PARALLEL QUERY



# SQL Query blocks determining FULL TABLE SCAN Operations [FTS]

- WHERE clause is missing (no filter predicate);
- WHERE clause with NULL condition;
- WHERE clause with predicates against unindexed columns;
- WHERE clause with LIKE condition (patterns like '%...');
- WHERE clause with NOT equals condition;
- WHERE clause with BIF (built-in-functions) invalidating index;
- **OEP hints** to enable table scan (full):
  - ALL\_ROWS hint;
  - PARALLEL hint;
  - FULL(table\_name | alias) hint;

# FTS up to: **HIGH WATER MARK** indicator



# Table Access Tuning TECHNIQUES

Tuning Technique	Table access operation	Notes
KEEP buffer	FTS (in memory)	Row resequencing and de-fragmentation to HWS
PARTITION ACCESS	PARTITION RANGE SINGLE PARTITION RANGE ITERATOR	Row resequencing and partition
Change access mode: access by ROWID	TABLE ACCESS BY INDEX ROWID	
Change access mode: access by ROWID	TABLE ACCESS BY USER RANGE ROWID	
Enable PARALLEL mode	Enable PARALLEL mode: PX	Multi-block I/O parameter

# Practice C5\_P1

Steps	Notes	SQL Script
1. HWM	High Water Mark	C5_P1.1.TABLE_ACCESS_HWM_FTS.sql
2. FTS Use Cases	OEP Hints: + FULL + PARALLEL + NOCACHE + CARDINALITY	C5_P1.1.TABLE_ACCESS_HWM_FTS.sql
3. Tuning: KEEP buffer	STORAGE (BUFFER_POOL KEEP)	C5_P1.2.TABLE_SCAN_Tuning.sql C5_P1.3.Base_Schema_CREATE_TABLE_KEEP.sql
4. Tuning: PARTITION-ing	PARTITION SCAN	C5_P1.2.TABLE_SCAN_Tuning.sql C5_P1.4.Base_Schema_CREATE_TABLE_STORAGE_RELOCATE.sql

# Practice C5\_P1

Steps	Notes	SQL Script
5. Tuning: change access mode to INDEX access	TABLE ACCESS BY INDEX ROWID	C5_P1.2.TABLE_SCAN_Tuning.sql
6. Tuning: change access mode to ROWID access	TABLE ACCESS BY RANGE ROWID	C5_P1.2.TABLE_SCAN_Tuning.sql
7. Tuning: reset HWT by re-formatting table storage	ALTER TABLE MOVE TABLESPACE DBMS_REDEFINITION. START_REDEF_TABLE	C5_P1.3.Base_Schema_CREATE_TABLE_STORAGE_RELOCATE.sql C5_P1.5.Base_Schema_ALTER_TABLE_STORAGE_RELOCATE.sql

## 2.2 INDEX-based optimization

- Index Design Process and Guidelines.
- Index-based OPERATIONS (from explain plan):
  - B\*Tree vs. Bitmap Specific operations
- Index-based tuning TECHNIQUES.

# Standard INDEX Design FACTORS to optimize SELECT Queries

- **Main factors** (and their relative importance) that **affect index design**:
  - **(1\*) filter predicates** (within WHERE clause) could reduce the number of scanned index segments (index slices):
    - through those filter-based columns evaluated as *matching columns*;
  - **(2\*) ordering criteria** (ORDER BY clause) included in the index-key structure could remove sorting operations executed after the table data access operations;
  - **(3\*) columns determining the resultset structure of data query** (SELECT column clause) included in the index-key structure could remove all table access operations needed:
    - “fat” indexing consequences need to be evaluated.

# Standard INDEX Design **PROCESS** to optimize SELECT Queries

- *1. Analyse WHERE predicates:*
  - 1.1 To choose the columns from *equals* predicates  
[col\_name=:value\_expression] to be added as first-key index column(s)  
(choosing matching columns);
  - 1.2 To choose the most selective predicate of *range* type [*>*, *<*, *>=*, *<=*,  
BETWEEN] to be added as *next-key column(s)* in the index structure.
- *2. Choose columns from ORDER BY clause* (taking in consideration also the DESCENDING sub-clause) to complete the index-key structure.
- *3. Add SELECT projection columns* to finalize the index-key structure, without some mandatory order.



## Other factors that could affect index design

- **Filter/Selection operators** from complex predicates like **range-type predicates** (containing operators like  $>$ ,  $<$ ,  $\geq$ ,  $\leq$ , or  $>|\geq\dots<|\leq$  combinations, or BETWEEN) :
  - could conflict with sorting criteria:
    - column order from ORDER BY not matching with the column priority derived from WHERE-predicates;
  - AND operator that link range-predicates could multiply index-slice scanning operations (multiple - index range scans);
  - OR operator that link range-predicates could produce less efficient multiple-index-access operations or even full-index-access operations.

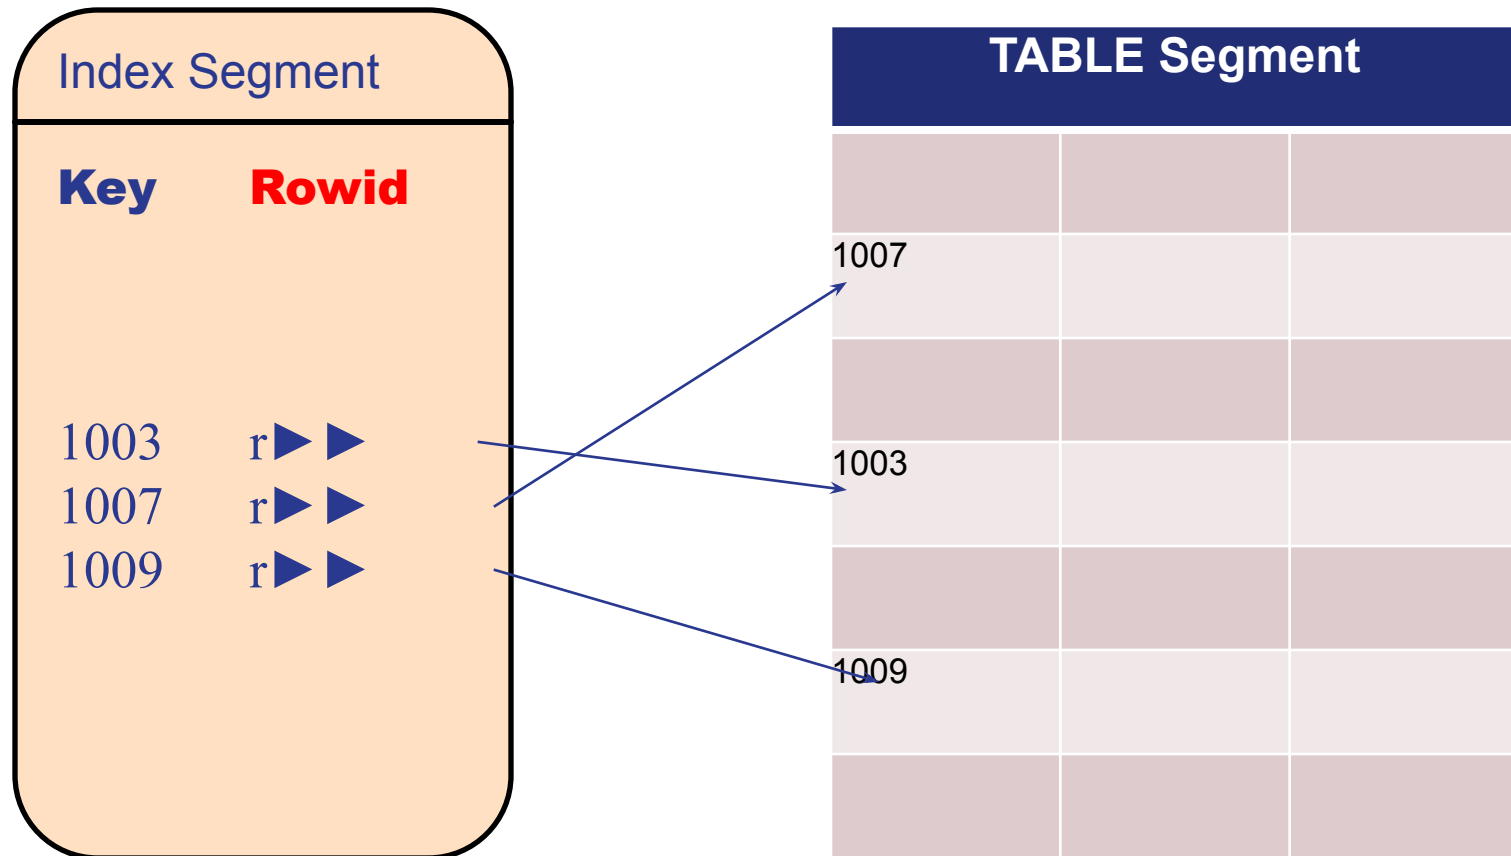
## Other factors that could affect index design

- Selection/Filter operators from complex predicates: such as **LIKE** with:
  - '%XX' search-pattern could disable a potential index scan or could determine a full index scan (instead of slice/range scan),
  - but 'XX%' search-pattern could determine a slice index scan.
- Selection/Filter operators from complex predicates such as **IN(multiple-value-list)**:
  - could determine multiple - index range scan operations;
  - if there are several *multiple - index range scan* predicates the most selective one could be chosen to determine the first column of the index-key structure.

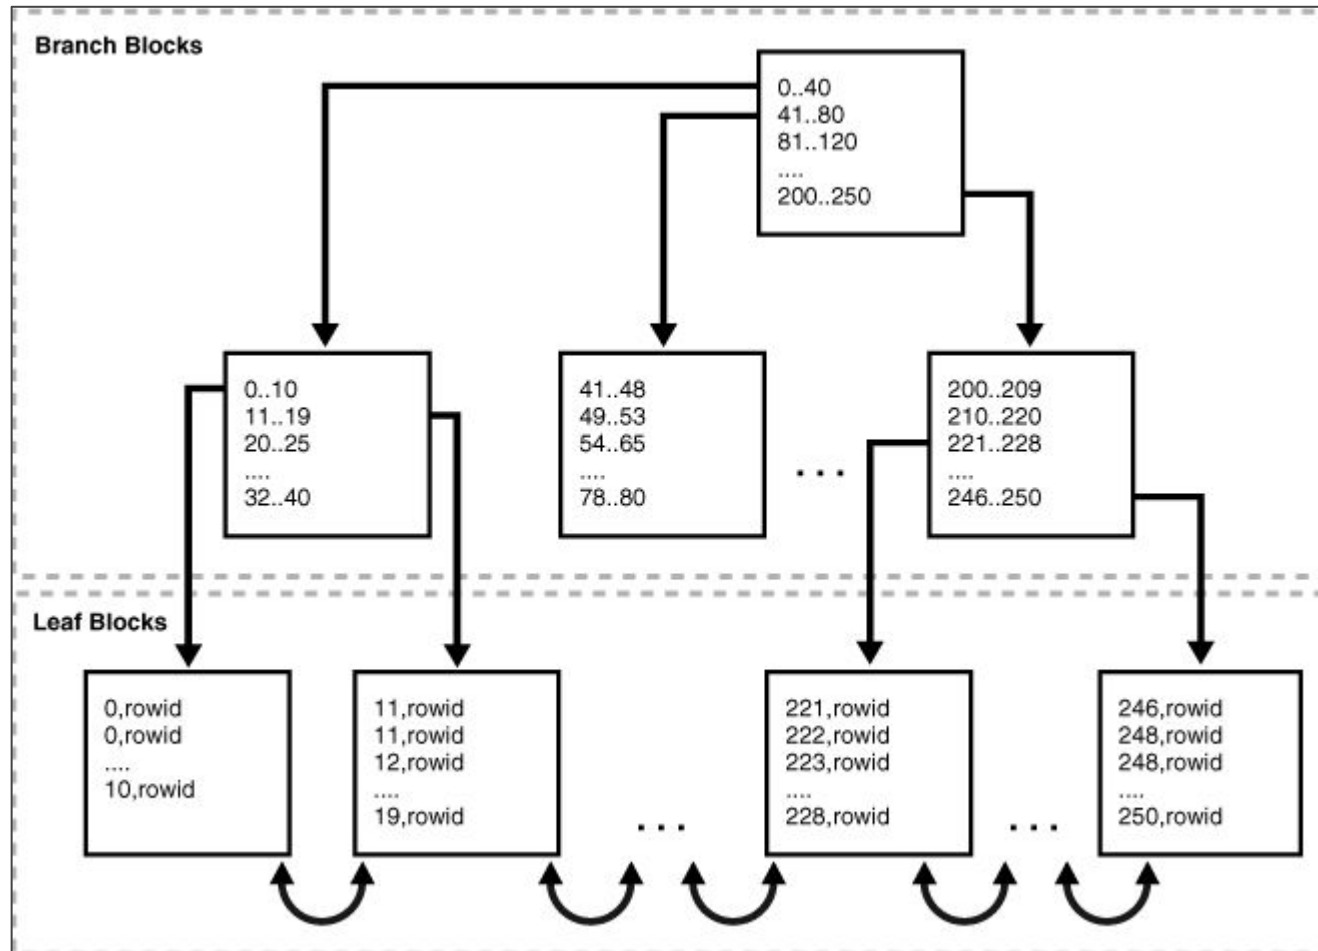
## Other factors that could affect index design

- **Not indexable** predicates (not “matching predicates”), e.g.:
  - [Col\_a] CONCAT [Col\_b] > [val];
  - [Col\_a] **NOT** BETWEEN [val\_1] AND [val\_2];
  - [Col\_a] LIKE [%val%];
- **Nested-JOIN predicates (and the ON clause):**
  - **Nested-JOIN** algorithm assumes scanning of the *driving table* (outer table) so that the probe table (inner table) will be iteratively scanned for each source-row from driving table;
  - some **WHERE** predicates on driving table could reduce the iterative scan operations on probe table, thus the indexed predicates will (indirectly) optimize (through a slice/range index scan) the whole NL process;
  - the actual **JOIN-ON** predicate analysis could suggest some indexing criteria for probe table (e.g. foreign-key based predicates) so that NL to be more efficiently probed not on actual table but on some index defined on the inner table.

# INDEX ROLE




# ORACLE B-TREE



# B\*TREE INDEX-based Operations

- **(preceding TABLE ACCESS BY INDEX)**
- B-tree Index Access
  - INDEX FULL SCAN
    - MIN|MAX
    - DESCENDING
  - INDEX RANGE SCAN
    - MIN|MAX
    - DESCENDING
  - INDEX SKIP SCAN
    - DESCENDING
  - INDEX UNIQUE SCAN
  - INDEX FAST FULL SCAN
  - INDEX JOIN

# ORACLE BITMAP



# BITMAP Index-based Operations

- BITMAP Index Access:
  - BITMAP INDEX SINGLE VALUE
  - BITMAP INDEX RANGE VALUE
  - BITMAP INDEX FULL SCAN
  - BITMAP INDEX FAST FULL SCAN
- Other BITMAP Index operations:
  - BITMAP CONVERSION
  - BITMAP OR
  - BITMAP AND
  - BITMAP MINUS
  - BITMAP MERGE
- *index\_combine* hint is used to force a bitmap access path for the table - if multiple bitmap indexes exist.



# Index-based Access optimization Techniques

SQL operation	Index-based Tuning Technique	Notes
SQL predicate: <ul style="list-style-type: none"><li>- equals predicates;</li><li>- range predicates;</li><li>- LIKE predicates;</li><li>- FBI predicates;</li></ul>	CREATE [B*TREE] INDEX CREATE BITMAP INDEX	LIKE predicate invalidates <a href="#">index usage</a> NULL predicate invalidates index usage BIF ON expressions or VIRTUAL columns
ORDER BY	CREATE [B*TREE] INDEX CREATE CLUSTER	Avoid SORT ORDER BY ( <a href="#">in memory</a> ) operations
SELECT-Projection	CREATE [B*TREE] INDEX	Cumulative: predicate + order_by_criteria + projection_columns
JOIN	CREATE B*TREE INDEX CREATE BITMAP INDEX	

# Practice C5\_P2

Steps	Notes	SQL Script
Gather Index Stats	DBMS_STATS. GATHER_INDEX_STATS	C5_P2.1.INDEX_ACCESS_Operations_and_Stats.sql
B*TREE Index Operations	FULL SCAN RANGE SCAN UNIQUE SCAN FAST FULL SCAN	C5_P2.1.INDEX_ACCESS_Operations_and_Stats.sql
BITMAP Index Operations	BITMAP AND BITMAP MINUS BITMAP MERGE	C5_P2.1.INDEX_ACCESS_Operations_and_Stats.sql

# Practice C5\_P2

Steps	Notes	SQL Script
INDEX and predicates	Equals predicates Range predicates LIKE predicates BIF predicates	C5_P2.2.INDEX_Based_Tuning.sql
INDEX-based tuning for sorting	ORDER BY	C5_P2.2.INDEX_Based_Tuning.sql
INDEX-based tuning for projection	SELECT Project	C5_P2.2.INDEX_Based_Tuning.sql
INDEX-based tuning for JOINS	NL and SORT-MERGE	C5_P2.2.INDEX_Based_Tuning.sql
INDEX-based tuning for NULL predicates	B-TREE and NULLs BITMAP and NULLs	C5_P2.2.INDEX_Based_Tuning.sql
Parallel INDEX SCAN	<code>/*+ parallel_index(t i)*/</code>	C5_P2.2.INDEX_Based_Tuning.sql

# References

- Craig S. Mullins, *Database Administration The Complete Guide to DBA Practices and Procedures* Second Edition, Addison-Wesley, 2013
- Tony Hasler, *Expert Oracle SQL Optimization, Deployment, and Statistics*, Apress, 2014
- Christian Antognini, *Troubleshooting Oracle Performance*, Apress, 2014
- Donald Burleson, *Oracle High-Performance SQL Tuning* 1st Edition, McGraw-Hill Education; 1 edition (August 17, 2001)

# References

- Lahdenmaki, Tapio, Leach, Michael, *Relational database index design and optimizers: DB2, Oracle, SQL server et al*, John Wiley & Sons, 2005
- Harrison, Guy, *Oracle performance survival guide: a systematic approach to database optimization*, Prentice Hall, 2009
- Allen, Grant, Bryla, Bob, Kuhn, Darl, *Oracle SQL Recipes: A Problem-Solution Approach*, Apress, 2009
- Caffrey, Mellanie et.al. *Expert Oracle Practices: Oracle Database Administration from the Oak Table*, Apress, 2010

# REFERENCES

- Oracle OTN Docs [[link 11gR2](#)]
- Oracle-Base Docs [[link 11gR2](#)]
- Toad World Docs [[link](#)]
- <http://www.orafaq.com/tuningguide/hints.html>