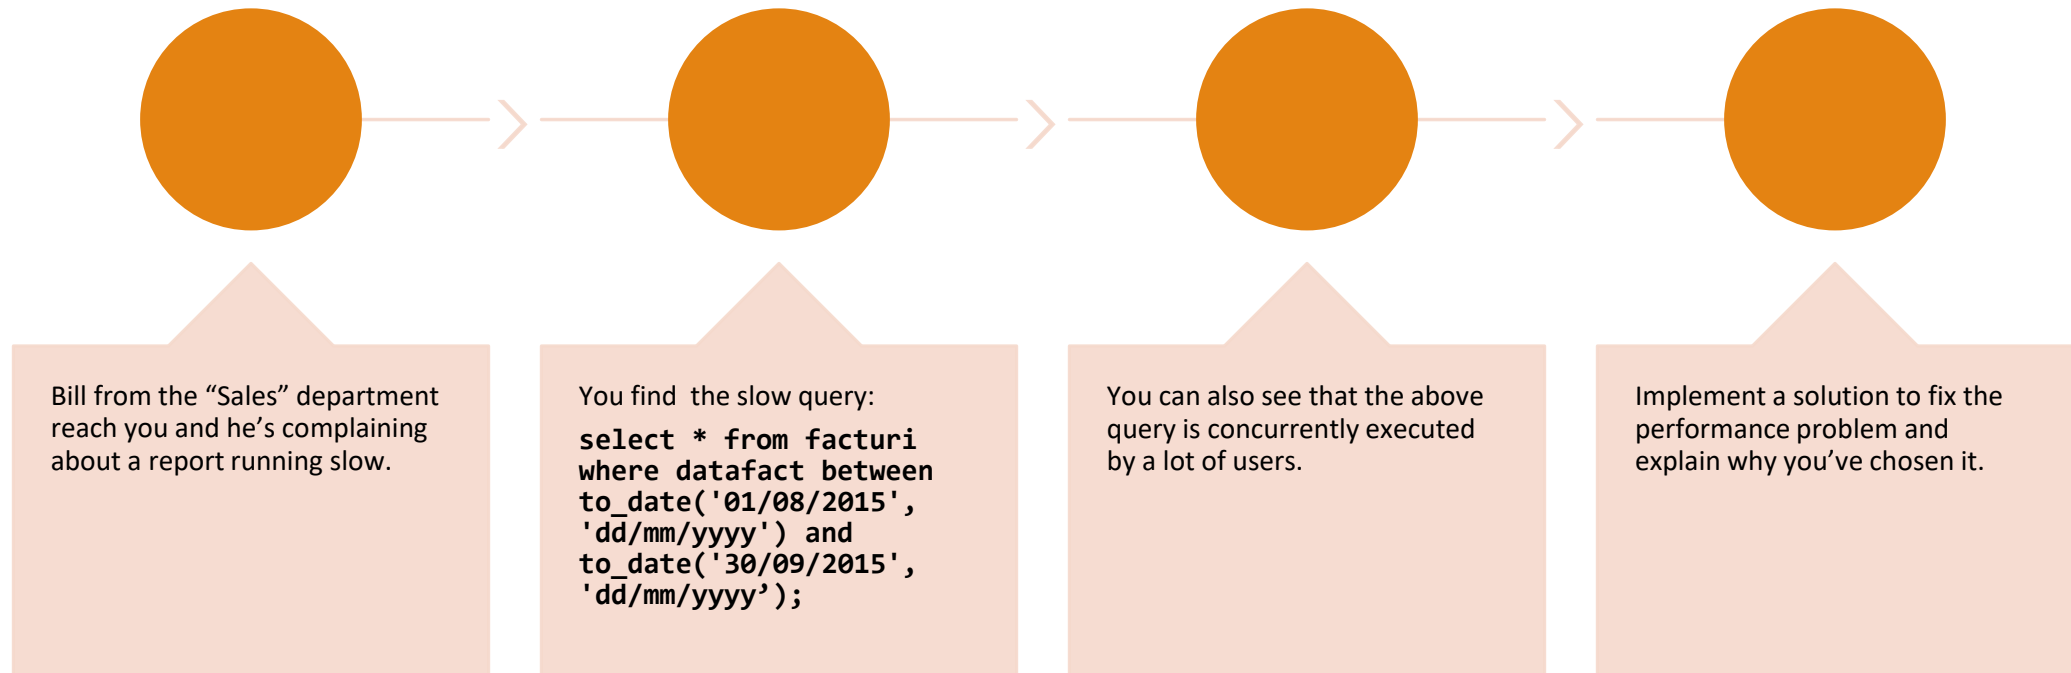


Oracle Performance Tuning

WEEK 8

Challenge (VANZARI schema)

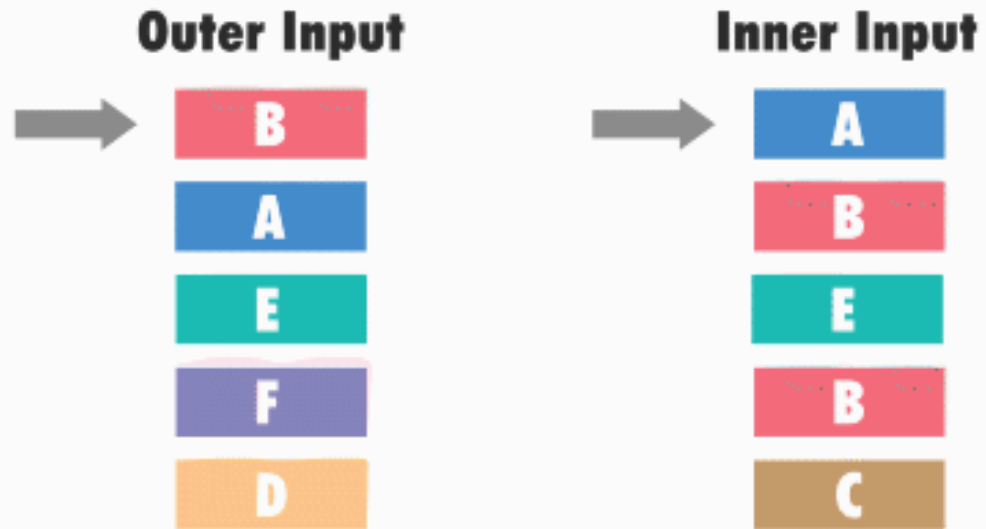




Optimizing JOINS

HOW MANY JOIN TYPES DO YOU
KNOW?

Nested Loops Join



SQL
BERT
bertwagner.com

The Nested
Loop Join

?!


```
SELECT /*+ use_nl(c cp) gather_plan_statistics*/ c.codcl, c.denc1, cp.loc
FROM coduri_postale cp INNER JOIN clienti c ON CP.CodPost = C.CodPost;
```

Query Result x Autotrace x

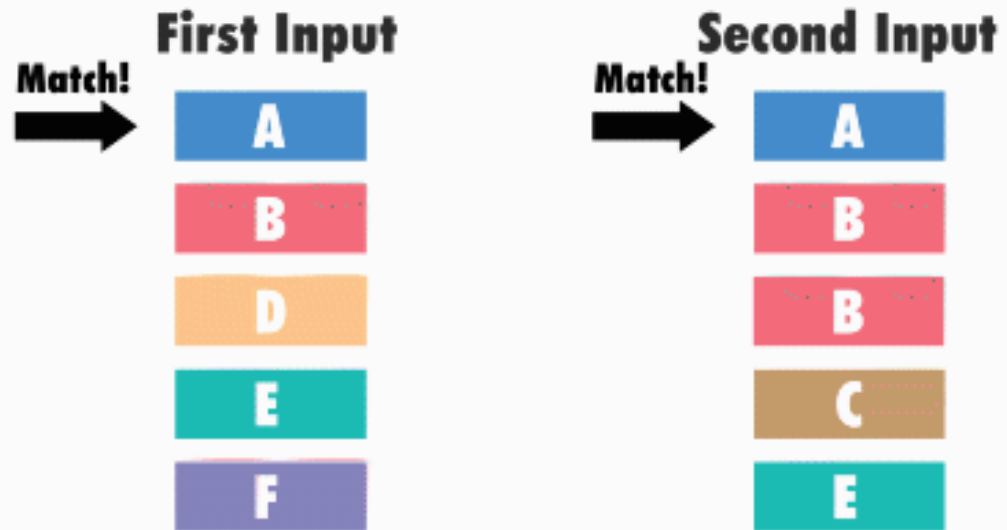
SQL HotSpot | 0.803 seconds

| OPERATION | OBJECT_NAME | CARDINALITY | COST |
|----------------------|----------------|-------------|------|
| SELECT STATEMENT | | | |
| NESTED LOOPS | | 1000 | |
| TABLE ACCESS (FULL) | CODURI_POSTALE | 10 | |
| TABLE ACCESS (FULL) | CLIENTI | 100 | |
| Filter Predicates | | | |
| CP.CODPOST=C.CODPOST | | | |

| V\$STATNAME Name | V\$MYSTAT Value |
|-----------------------------|-----------------|
| cluster key scans | 9 |
| consistent gets | 233 |
| consistent gets examination | 71 |

NL Example

Merge Join



SQL
BERT
bertwagner.com

Sort Merge Join

| <pre> SELECT /*+ gather_plan_statistics use_merge(c f l p) */ c.dencl, f.nrfact, f.datafact, p.denpr, l.cantitate FROM clienti c INNER JOIN facturi f ON c.codcl = f.codcl INNER JOIN liniifact l ON f.nrfact = l.nrfact INNER JOIN produse p ON l.codpr = p.codpr ORDER BY 1, 2, 3; </pre> | | | | |
|---|-------------|-------------|-------|--|
| Query Result x Autotrace x | | | | |
| SQL HotSpot 2.565 seconds | | | | |
| OPERATION | OBJECT_NAME | CARDINALITY | COST | |
| SELECT STATEMENT | | | 10484 | |
| SORT (ORDER BY) | | 151393 | 10484 | |
| MERGE JOIN | | 151393 | 5961 | |
| SORT (JOIN) | | 151393 | 5957 | |
| MERGE JOIN | | 151393 | 2378 | |
| SORT (JOIN) | | 26280 | 685 | |
| MERGE JOIN | | 26280 | 276 | |
| SORT (JOIN) | | 1000 | 6 | |
| TABLE ACCESS (F CLIENTI) | | 1000 | 5 | |
| SORT (JOIN) | | 26280 | 270 | |
| Access Predicates | | | | |
| Filter Predicates | | | | |
| TABLE ACCESS (F FACTURI) | | 26280 | 22 | |
| Filter Predicates | | | | |
| SORT (JOIN) | | 151393 | 1693 | |
| Access Predicates | | | | |
| Filter Predicates | | | | |
| TABLE ACCESS (FULL) LINIIFACT | | 151393 | 171 | |
| Filter Predicates | | | | |
| SORT (JOIN) | | 1000 | 4 | |
| Access Predicates | | | | |
| Filter Predicates | | | | |
| TABLE ACCESS (FULL) PRODUSE | | 1000 | 3 | |

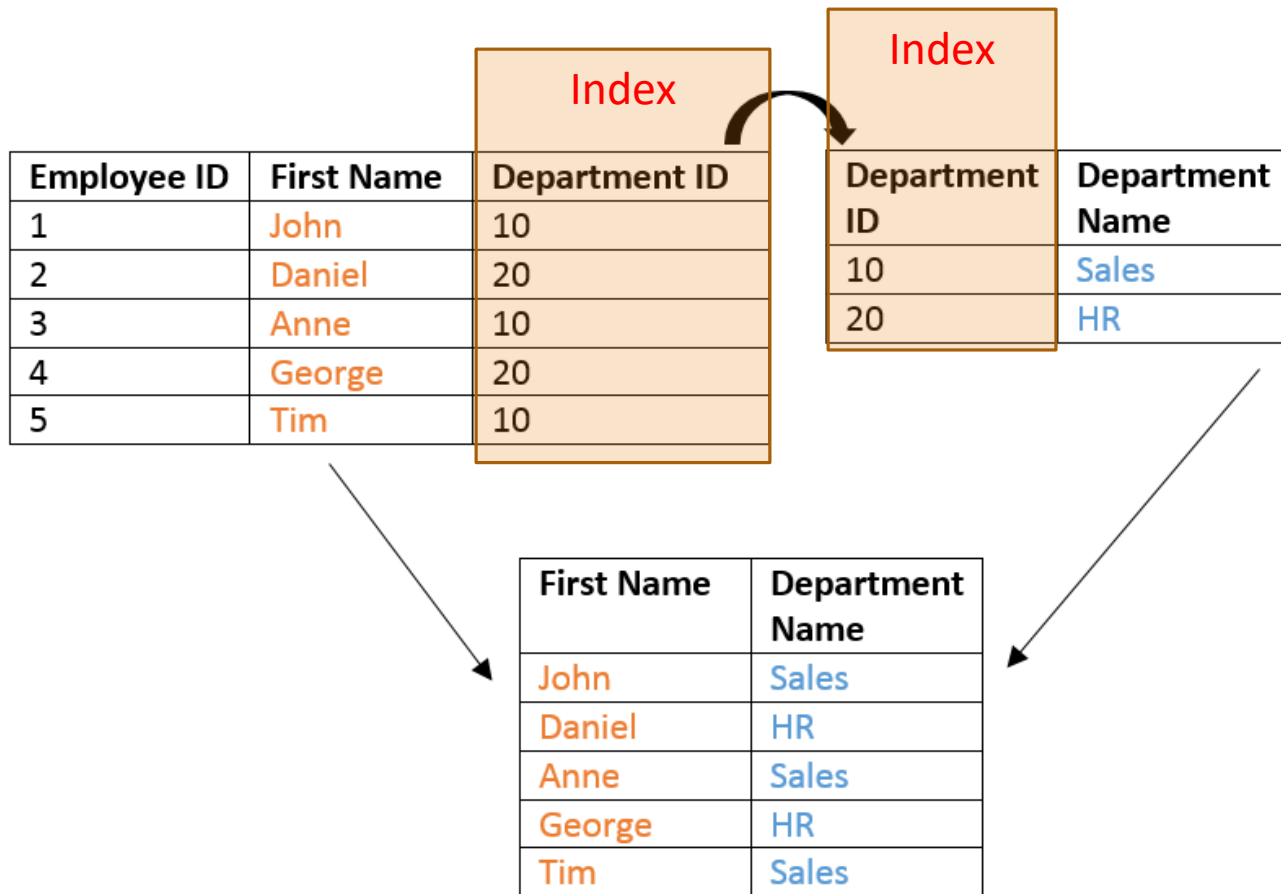
Sort MJ Example

divide and conquer

HASH Joins

| <pre> SELECT /*+ gather_plan_statistics use_hash(c f l p) */ c.dencl, f.nrfact, f.datafact, p.denpr, l.cantitate FROM clienti c INNER JOIN facturi f ON c.codcl = f.codcl INNER JOIN liniifact l ON f.nrfact = l.nrfact INNER JOIN produse p ON l.codpr = p.codpr; </pre> | | | |
|---|-------------|-------------|------|
| Query Result x Autotrace x | | | |
| SQL HotSpot 0.587 seconds | | | |
| OPERATION | OBJECT_NAME | CARDINALITY | COST |
| SELECT STATEMENT | | | 627 |
| HASH JOIN | | 151393 | 627 |
| Access Predicates L.CODPR=P.CODPR | | | |
| TABLE ACCESS (FULL) | PRODUSE | 1000 | 3 |
| HASH JOIN | | 151393 | 624 |
| Access Predicates C.CODCL=F.CODCL | | | |
| TABLE ACCESS (FULL) | CLIENTI | 1000 | 5 |
| HASH JOIN | | 151393 | 618 |
| Access Predicates F.NRFACT=L.NRFACT | | | |
| TABLE ACCESS (FULL) | FACTURI | 26280 | 22 |
| Filter Predicates F.CODCL>1000 | | | |
| TABLE ACCESS (FULL) | LINIIFACT | 151393 | 171 |
| Filter Predicates L.CODPR>0 | | | |

HJ Example



Indexing Joining Columns

Bitmap Join Indexes

EMPLOYEES table

| Employee_id | Name | Department_id |
|-------------|---------|---------------|
| 1 | Larry | 1 |
| 2 | Ken | 1 |
| 3 | Ray | 2 |
| 4 | Married | 3 |
| 5 | Married | 3 |
| 6 | Single | 2 |
| 7 | Married | 2 |
| | | |

DEPARTMENTS table

| Department_id | Name |
|---------------|-----------|
| 1 | Sales |
| 2 | R&D |
| 3 | Marketing |

| Sales | R&D | Marketing |
|-------|-------|-----------|
| 1 | 0 | 1 |
| 1 | 0 | 0 |
| 0 | 1 | 0 |
| 0 | 0 | 1 |
| 0 | 0 | 1 |
| 0 | 1 | 0 |
| 0 | 1 | 0 |
| | | |

Bitmap Join Index

```
CREATE BITMAP INDEX bi_1 ON employees(d.department_name)
FROM employees e, departments d WHERE d.department_id=e.department_id
```

Bitmap Join Example (part 1)

The Query:

```
SELECT COUNT(*) FROM CLIENTI C, FACTURI F
WHERE C.CODCL=F.CODCL AND C.DENCL='Client 1019';
```

| OPERATION | OBJECT_NAME | CARDINALITY | COST |
|--|-------------|-------------|------|
| SELECT STATEMENT | | | 27 |
| SORT (AGGREGATE) | | 1 | |
| HASH JOIN | | 29 | 27 |
| Access Predicates C.CODCL=F.CODCL | | | |
| TABLE ACCESS (FULL) | CLIENTI | 1 | 5 |
| Filter Predicates C.DENCL='Client 1019' | | | |
| TABLE ACCESS (FULL) | FACTURI | 26280 | 22 |
| Filter Predicates F.CODCL>1000 | | | |

| | |
|------------------|-----------------|
| consistent gets | |
| V\$STATNAME Name | V\$MYSTAT Value |
| consistent gets | 468 |

Bitmap Join Example (part 2)

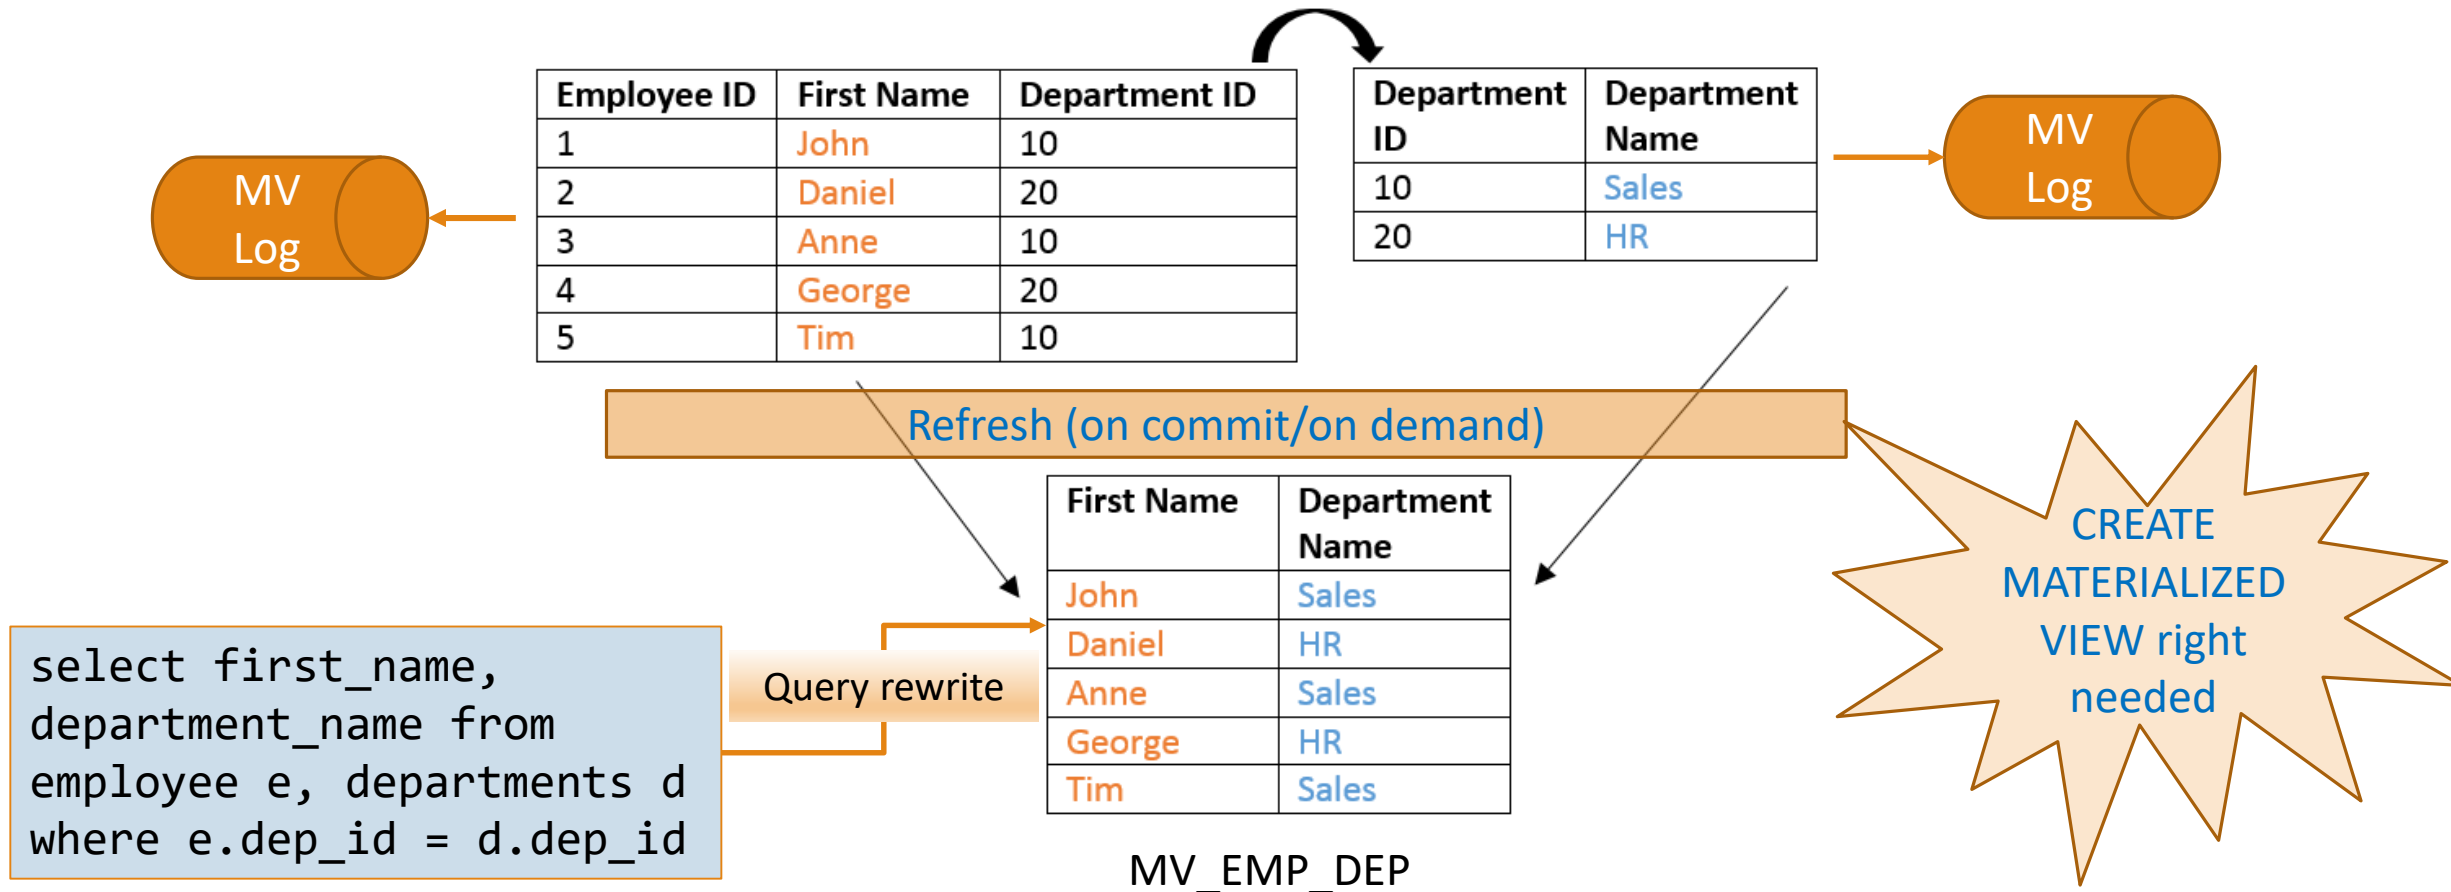
```
CREATE BITMAP INDEX BIDX_FACTURI_DENCL ON  
FACTURI(C.DENCL) FROM FACTURI F, CLIENT C  
WHERE C.CODCL = F.CODCL;
```

| OPERATION | OBJECT_NAME | CARDINALITY | COST |
|--|-------------|-------------|------|
| SELECT STATEMENT | | | 1 |
| SORT (AGGREGATE) | | 1 | |
| TABLE ACCESS (BY INDEX ROWID BATC FACTURI) | | 27 | 1 |
| Filter Predicates | | | |
| F.CODCL>1000 | | | |
| BITMAP CONVERSION (TO ROWIDS) | | | |
| BITMAP INDEX (SINGLE VALUE) BIDX_FACTURI_DENCL | | | |
| Access Predicates | | | |
| F.SYS_NC00008\$='Client' | | | |

cons|

| V\$STATNAME Name | V\$MYSTAT Value |
|------------------|-----------------|
| consistent gets | 20 |

Materialized Views



There is an application which is using data from “VANZARI” schema

Some users are mad because one of their frequently used report is running slow

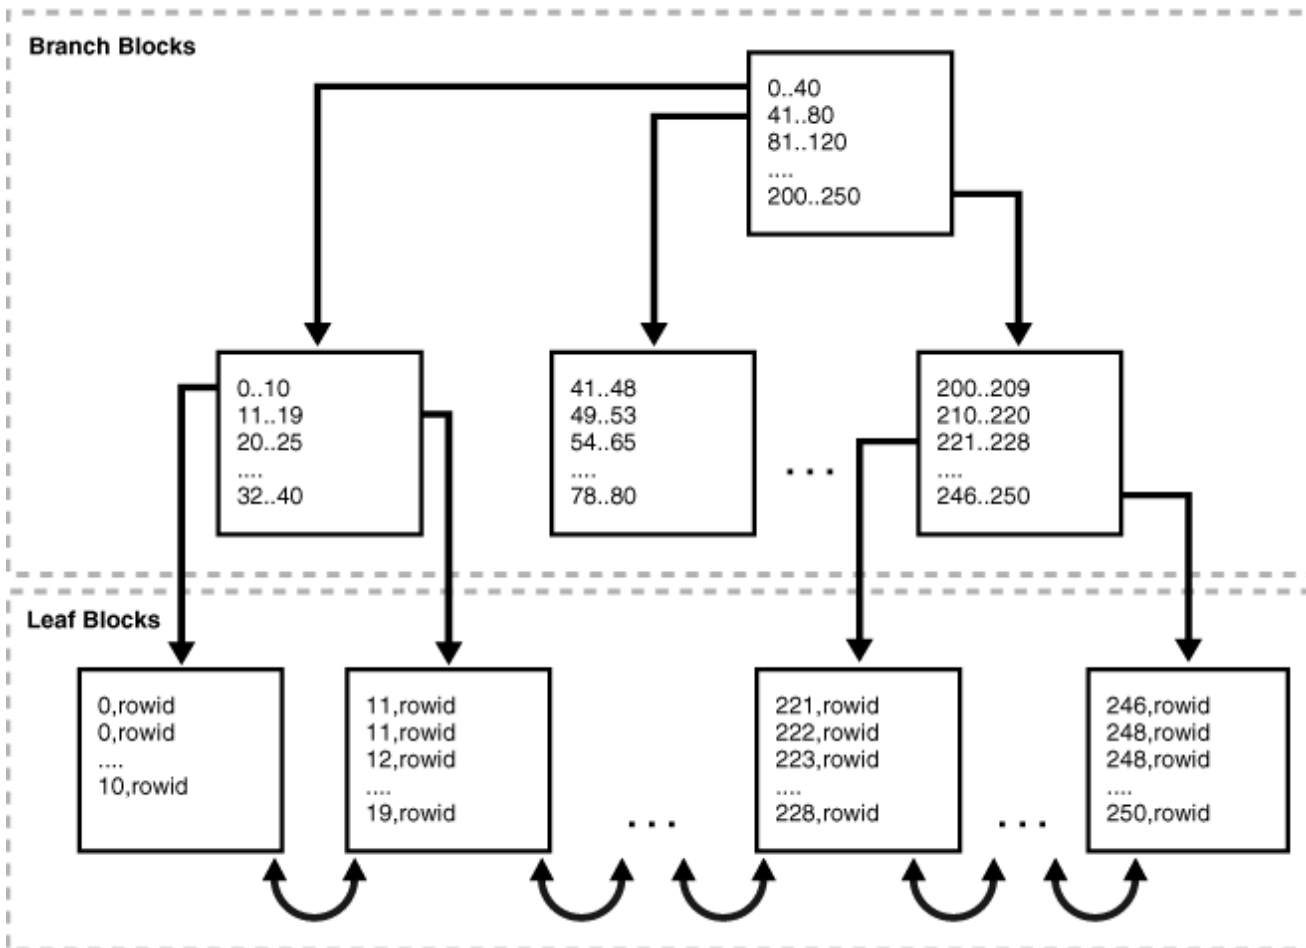
You find out that the following query is behind that report: `SELECT F.NRFACT, L.CODPR, L.CANTITATE FROM FACTURI F, LINIIFACT L WHERE F.NRFACT=L.NRFACT AND F.DATAFACT < TO_DATE('01/09/2015', 'DD/MM/YYYY')`

Suggest and implement an improvement for the above query

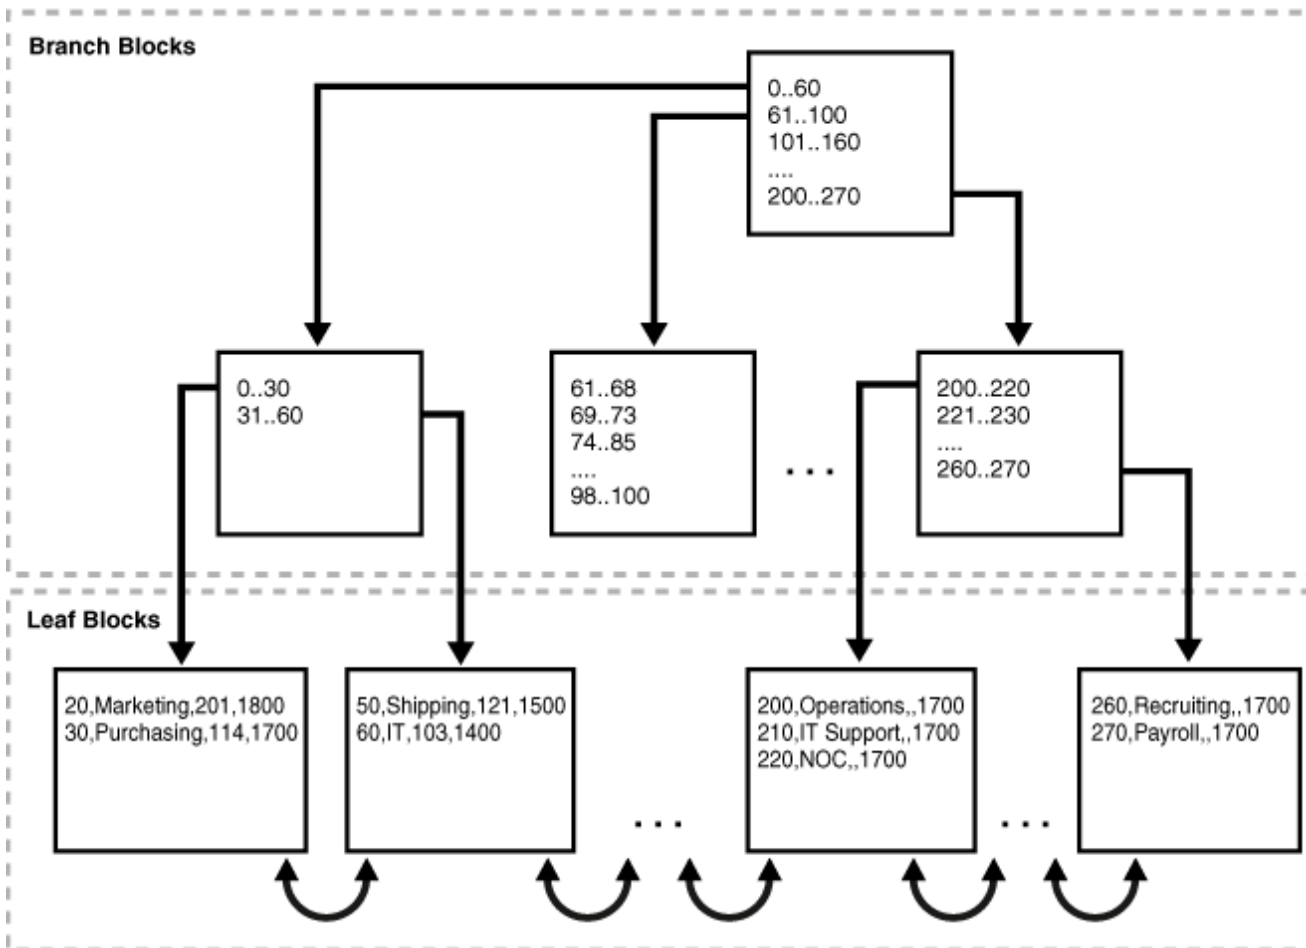
How can you tell that the fix you suggested did the job?

Hands-on Practice

Other SQL Tuning Techniques



Do you
remember
indexes?



Tabele IOT

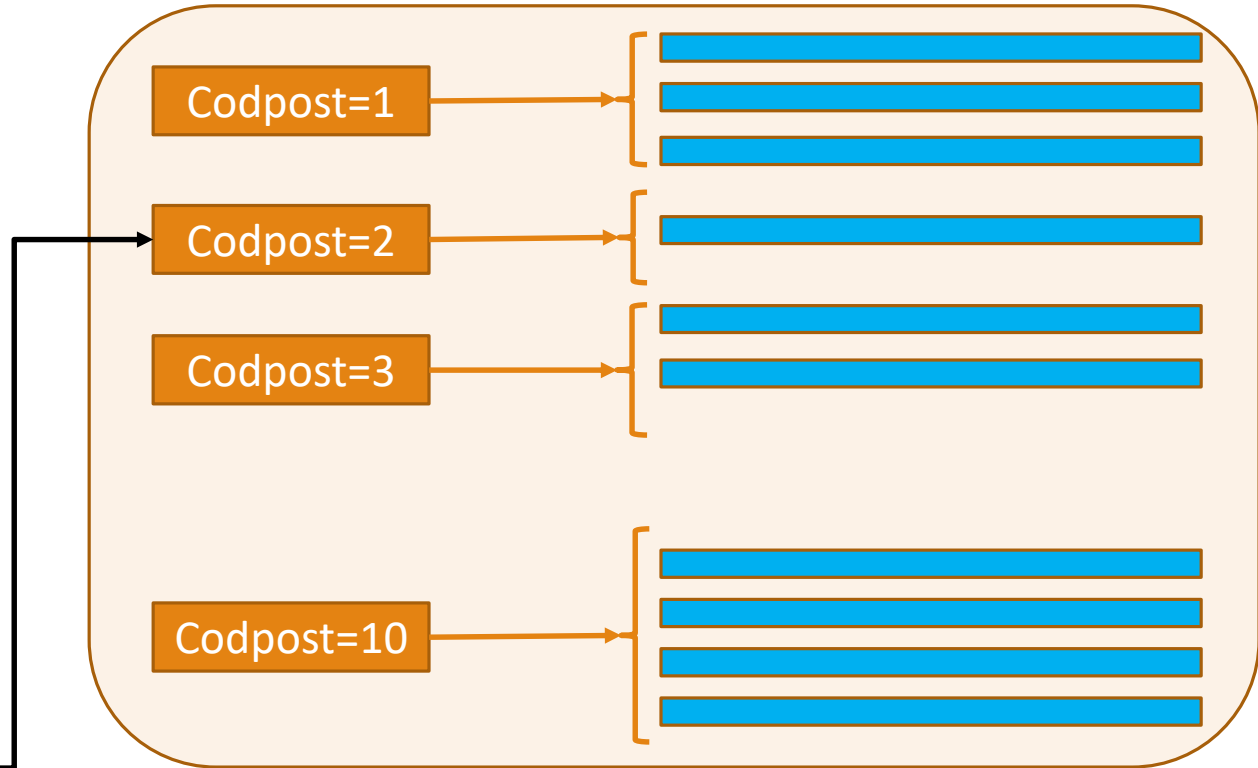
| Heap-Organized Table | Index-Organized Table |
|---|--|
| The rowid uniquely identifies a row. Primary key constraint may optionally be defined. | Primary key uniquely identifies a row. Primary key constraint must be defined. |
| Physical rowid in ROWID pseudocolumn allows building secondary indexes. | Logical rowid in ROWID pseudocolumn allows building secondary indexes. |
| Individual rows may be accessed directly by rowid. | Access to individual rows may be achieved indirectly by primary key. |
| Sequential full table scan returns all rows in some order. | A full index scan or fast full index scan returns all rows in some order. |
| Can be stored in a table cluster with other tables. | Cannot be stored in a table cluster. |
| Can contain a column of the LONG data type and columns of LOB data types. | Can contain LOB columns but not LONG columns. |
| Can contain virtual columns (only relational heap tables are supported). | Cannot contain virtual columns. |

Regular tables vs. IOTs

Hash Clusters

```
create cluster coduri_postale_hash (  
  codpost char(6)  
)  
hashkeys 10  
size 8192;  
  
create table clienti_cluster (  
  codcl number(6),  
  denc1 varchar2(30),  
  codfiscal char(9),  
  adresa varchar2(40),  
  codpost char(6),  
  telefon varchar2(10)  
) cluster coduri_postale_hash(codpost);
```

```
select * from  
clienti_cluster where  
codpost=2
```



That's all folks!

THANK YOU AND SEE YOU NEXT WEEK...